

Live coding environment for a data exploration language

ANONYMOUS AUTHOR(S)

A growing amount of code is written to explore and analyze data, but the way data analysts work with code differs from the way software engineers do. First, many data scientists write code that introduces few or no abstractions. Second, data exploration is an interactive process and it blurs the traditional distinction between development-time and run-time. This poses an interesting challenge for programming language research, because it allows us to revisit a number of traditional assumptions about programming languages and focus our attention on a new kind of programming tools that data analysts need.

In this paper, we present *data exploration calculus*, a formally tractable language that captures the structure of simple data analyses as done, for example, by journalists exploring open government data. We implement a programming environment for data exploration that evaluates code instantaneously during editing and shows previews of the results, while allowing the user to modify code in an unrestricted way in a text editor. Supporting a text editor is tricky as any edit can change the structure of code and fully recomputing program after every change would be too expensive. We present a new technique that allows evaluating and type-checking code while it is being written while reusing results of previous runs. We formalize the technique using the data exploration calculus, prove that it is correct and specify when previous results are reused. Finally, we also illustrate the practicality of our approach using empirical evaluation and a case study.

As data analysis becomes ever more important use of programming, research on programming languages and tools needs to consider new kinds of code that people need to write and new kinds of tools that can support them. The present paper is one step in this important direction.

1 INTRODUCTION

One of the aspects that make spreadsheets easier to use than programming tools is their liveness. When you change a value in a cell in Excel, the whole spreadsheet updates instantly and you immediately see new results, without having to explicitly trigger re-computation and without having to wait for an extensive period of time.

Increasing number of programming environments aim to provide the same live development experience for standard programming languages, but doing this is not easy. Fully recomputing the whole program after every keystroke is inefficient and calculating how a change in the source code changes the result is extremely hard when the text editor allows arbitrary changes. Consider the following snippet that gets the release years of 10 most expensive movies from a data set `movies`:

```
let top = movies.sortBy( $\lambda x \rightarrow x.getBudget()$ )
               .take(10).map( $\lambda x \rightarrow x.getReleased().format("yyyy")$ )
```

A live coding environment computes and shows the list of years, but then the programmer modifies the code by making the constant `10` a variable and changing the date format to see the full date:

```
let count = 10
let top = movies.sortBy( $\lambda x \rightarrow x.getBudget()$ )
               .take(count).map( $\lambda x \rightarrow x.getReleased().format("dd-mm-yyyy")$ )
```

Ideally, the live coding environment should understand the change, reuse a cached result of the first two transformations (sorting and taking first 10 elements) and only evaluate the `map` operation to differently format the release dates of already computed top 10 movies. We present such a live coding environment, describe it formally and empirically evaluate it. We discuss related work in Section 8, but we review the most important directions next, in order to situate our contributions.

1.1 Related work

We focus on simple data exploration as done, for example, by journalists [Gray et al. 2012] and we aim to bring the experience of using transparent and reproducible programming tools closer to that of using spreadsheets. When a user edits a cell in a spreadsheet, the results automatically update. Recomputation in spreadsheets is well-studied, e.g. by Sestoft [2012], but it is simpler than in our case, because editing is limited to individual cells.

Many visual data exploration tools support interactivity [Crotty et al. 2015; Hellerstein et al. 1999; Wesley et al. 2011] and some, can even export the workflow as a script [Kandel et al. 2011], but data analysts who prefer working with code typically resort to notebook systems such as Jupyter or R Markdown [Allaire et al. 2016; Kluyver et al. 2016]. Those are text-based, but have a limited model of recomputation. Users have to structure code in cells and manually reevaluate cells after a change. Those systems also do not track dependencies between cells, which can lead to inconsistencies; an issue that has been addressed by e.g. Koop and Patel [2017] or Petricek et al. [2018].

We aim to keep the transparency and reproducibility of notebooks, but add a correct and efficient instantaneous evaluation mechanism. Existing text-based editors that provide instant, such as Lighttable [Granger 2012] and Chrome DevTools, often work only in certain situations or require full recomputation. A more principled approach has been used by systems based on structured editing [Lotem and Chuchem 2018; Omar et al. 2019], which benefit from the fact that code can only be modified via known operations with known effect on the computation graph (“extract variable” has no effect on the result; “change constant value” forces recomputation of subsequent code). However, we aim to cater for the many users who prefer to edit programs as free-form text.

The technical aspects of our technique for providing instantaneous feedback are related to the work on incremental computation [Acar 2005; Hammer et al. 2015]. This focuses on cases where *data* changes but *program* stays the same, while our focus is on cases where *program* changes but *data* stays. Yet, the technical aspects are related and, as noted in Section 8, language that supports incremental computation could be used to implement the kind of system presented in this paper.

1.2 Contributions

We present the design and implementation of a live coding environment for a data exploration language that provides a correct and efficient instantaneous feedback, yet, is integrated into an ordinary text editor. Our most important contributions are:

- We introduce *data exploration calculus* (Section 3), a small formally tractable programming language for data exploration. The language is motivated by our review of how data analysts work (Section 2) and captures the idea that scripts are often lists of commands with few or no abstractions with most logic provided by external libraries.
- Implementing a live coding environment requires a different way of thinking than the one presented in classic compiler literature. We capture the essence of the new perspective (Section 4) and use it to build our *live preview* mechanism (Section 5) evaluates code instantaneously during editing.
- We prove that our live preview mechanism is correct (Section 6.1) and that it reuses previously evaluated values when possible (Section 6.2). We illustrate the practicality of the mechanism using an empirical evaluation (Section 6.3) and a case study (Section 6.4).
- We extend our core environment in two ways, First, we add instantaneous type checker that relies on the same core infrastructure as live previews and enables richer user experience through *type providers* (Section 7.1). Second, we add an abstraction mechanism that preserves the live coding experience (Section 7.2).

2 UNDERSTANDING HOW DATA SCIENTISTS WORK

Data scientists often use general-purpose programming languages such as Python, but the kind of code they write and the way they interact with the system is very different from how software engineers work [Guo 2013]. This paper focuses on simple data wrangling and data exploration as done, for example, by journalists analysing government datasets. This section illustrates how such data analyses look and provides a justification for the design of our data exploration calculus.

2.1 Data wrangling in Jupyter

Data analysts increasingly use notebook systems such as Jupyter, which make it possible to combine text, equations and code with results of running the code such as tables or visualizations. Notebooks blur the distinction between development and execution. Data analysts write small snippets of code, run them to see results immediately and then revise them.

Jupyter notebooks are used by a variety of users ranging from scientists who implement complex models of physical systems to journalists who load data, perform simple aggregations and create visualizations. Our initial focus is on simple use cases, such as those of journalists. Tooling that makes notebook systems more interactive and spreadsheet-like can encourage users to choose programming tools over spreadsheets and produce reproducible and transparent data analyses.

As an example, the analysis done by the Financial Times for an article on plastic waste [Blood 2018; Hook and Reed 2018] joins datasets from Eurostat, UN Comtrade and more, aggregates the data and builds a visualization comparing waste flows in 2017 and 2018. Figure 1 shows an extract from one notebook from the data analysis. The code has a number of remarkable properties:

UN Comtrade exports data

```
import pandas as pd
material = 'plastics' # 'plastics', 'paper'
```

Loading exports data

```
df_mat = pd.read_csv('{material}-2017.csv').fillna(0).sort_values(['country_name', 'period'])
df_mat.head()
```

	period	country_name	kg	country_code
0	2017-01-01	Algeria	43346.0	12
1	2017-03-01	Algeria	32800.0	12
2	2017-03-01	Antigua and Barbuda	17000.0	28

Join to country codes

```
# Set keep_default_na because the Namibia has ISO code NA
df_isos = pd.read_excel('iso.xlsx', keep_default_na=False).drop_duplicates('country_code')
df = df_mat.copy().merge(df_isos, 'left', 'country_code').rename({'iso2': 'country_code'}, axis=1)
df.head()
```

	period	country_name	kg	country_code
0	2017-01-01	Algeria	43346.0	DZ
1	2017-03-01	Algeria	32800.0	DZ
2	2017-03-01	Antigua and Barbuda	17000.0	AG

Fig. 1. An excerpt from a Jupyter notebook by Financial Times reporters, which loads trade data from the UN trade database (stored in a local file) and adds ISO country codes (from an Excel spreadsheet).

- There is no abstraction. The analysis uses lambda functions as arguments to library calls, but it does not define any top-level functions. Parameterization is done by defining a list of inputs and then having a for loop, or by defining a top-level variable and setting its value. In our example, material is set to "plastics" and a comment lists other options. This lets the analyst easily see and check the results of intermediate steps of the analysis.
- The code relies on external libraries. Our example uses Pandas [McKinney 2012], which provides operations for data wrangling such as `merge` to join datasets or `drop_duplicates` to delete rows with duplicate column values. Those libraries are external to the data analysis and may even be implemented in other languages – many Python libraries rely on C++.
- The code is structured as a sequence of commands. Some commands define a variable, either by loading data, or by transforming data loaded previously. Even in Python, data is often treated as immutable. Other commands produce an output that is displayed in the notebook.
- There are many corner cases, such as the fact that the `keep_default_na` parameter needs to be set to handle Namibia correctly. These are discovered interactively by writing and running a code snippet and examining the output, so providing a rapid feedback is essential.

Many Jupyter notebooks are more complex than the example discussed here. They might define helper functions or even structure code using object-oriented programming. However, this paper is motivated by our conviction that simple data analyses such as the one discussed here are frequent enough to pose an interesting and important domain for programming tools research.

2.2 Dot-driven data exploration in The Gamma

Data exploration of the kind discussed in the previous section has been the motivation for a recent work by Petricek [2017b] that develops a small scripting language called The Gamma. The language follows the structure discussed above. A script is a sequence of commands that can either define a variable or produce an output. The language does not support top-level function declarations and lambda functions can be used only as method arguments.

```
olympics
.filter data'.Games is'.Rio (2016)'.then
.group data'.by Athlete'
.count distinct Event'
.su
```

```
count distinct Gender
count distinct Gold
count distinct Medal
count distinct Silver
count distinct Sport
count distinct Team
count distinct Year
sum Bronze
sum Gold
sum Silver
sum Year
then
```

(a) The Gamma script to aggregate Olympic medals. The analyst selects Rio 2016 games and counts the number of distinct events per athlete. After typing `;`, a type provider offers further aggregation operations.

```
olympics
.filter data'.Games is'.Rio (2016)'.then
.group data'.by Athlete'
.count distinct Event'.sum Gold'.then
.sort data'.by Gold descending'.then
.paging.take(5)
```

Athlete	Event	Gold
Michael Phelps	6	5
Katie Ledecky	5	4
Simone Biles	5	4
Katinka Hosszu	4	3
Usain Bolt	3	3

(b) Live preview in The Gamma programming environment developed based on the theory presented in this paper. The table is produced as the data analyst types and updates on-the-fly to show result at the current cursor position.

Fig. 2. Data exploration in The Gamma – using a type provider (left) and with a live preview (right).

Given the limited expressiveness of The Gamma, libraries are implemented in other languages, such as JavaScript. The Gamma also supports type providers [Syme et al. 2013], which can be used to generate types, on the fly, for accessing and exploring external data sources. Type providers provide object types with members and The Gamma makes using those convenient by providing auto-complete when the user types the dot symbol (‘.’) to access a member.

As illustrated by Petricek [2017b], the combination of type providers and auto-complete on ‘.’ makes it possible to complete a large number of data exploration tasks through a very simple interaction of selecting operations from a list. An example in Figure 2a summarizes data on Olympic medals using a type provider. The names of identifiers such as ‘sum Bronze’ do not have any special meaning. They are merely members generated by the type provider, based on information about the data source. The type provider used in this example generates an object with members for different data transformations, such as ‘group data’, which return further objects with members for specifying transformation parameters, such as selecting the grouping key using ‘by Athlete’.

The Gamma language is richer, but the example in Figure 2a shows that non-trivial data exploration can be done using a very simple language. The assumptions about structure of code that are explicit in The Gamma are implicitly present in Python and R data analyses produced by journalists, economists and other users with other than programming background. When we refer to The Gamma in this paper, readers not familiar with it can assume that we focus on a subset of Python.

2.3 Live coding for data exploration

The practical implementation work that complements the theory presented in this paper implements a live coding environment for The Gamma and is discussed in Section 6.4 in detail. Earlier work by Petricek [2017b] discussed The Gamma language and the type provider mechanism, but the implementation provided only a standard text editor with auto-completion. Our work brings the experience of using The Gamma closer to that of spreadsheets by developing an editor that evaluates scripts instantaneously in background and displays the results as *live previews*.

An example of a live preview is shown in Figure 2b. As noted earlier, The Gamma program consists of a list of commands which are either expressions or let bindings. Our editor displays a live preview below the command that the user is currently editing. The preview shows the result of evaluating the expression or the value assigned to a bound variable. When the user changes the code, the preview is updated automatically, without any additional interaction with the user.

There is a number of guiding principles that our work follows. First, we allow the analyst to edit code in an unrestricted form in a text editor. Although structured editors provide an appealing alternative and make recomputation easier, we prefer the transparency and openness of plain text. Second, we focus on the scenario when code changes, but input does not. Rapid feedback allows the analyst to quickly adapt code to correctly handle corner cases that typical analysis involves. In contrast to work on incremental computation, we do not consider the case when data changes, although supporting interactive data exploration of streaming data is an interesting future direction.

3 DATA EXPLORATION CALCULUS

The *data exploration calculus* is a small, formally tractable language for data exploration. The calculus is not Turing-complete and it can only be used together with external libraries that define what objects are available and what the behaviour of their members is. This is sufficient to capture the simple data analyses discussed in Section 2. We define the calculus in this section and then use it as the basis for discussing our live preview mechanism in Section 4.

We initially present the data exploration calculus without a static type system as the live preview mechanism does not rely on types. We consider type checking later in Section 7.1 as the mechanism that is used for live previews can also be used for instantaneous type checking.

Programs, commands, terms, expressions and values

$$\begin{array}{lll} p = c_1; \dots; c_n & t = o \mid x & e = t \mid \lambda x \rightarrow e \\ c = \text{let } x = t \mid t & \mid t.m(e, \dots, e) & v = o \mid \lambda x \rightarrow e \end{array}$$

Evaluation contexts of expressions

$$\begin{array}{ll} C_e[-] = C_e[-].m(e_1, \dots, e_n) \mid o.m(v_1, \dots, v_m, C_e[-], e_1, \dots, e_n) \mid - \\ C_c[-] = \text{let } x = C_e[-] \mid C_e[-] \\ C_p[-] = o_1; \dots; o_k; C_c[-]; c_1; \dots; c_n \end{array}$$

Let elimination and member reduction

$$\begin{array}{ll} o_1; \dots; o_k; \text{let } x = o; c_1; \dots; c_n \rightsquigarrow & \text{(let)} \\ o_1; \dots; o_k; o; c_1[x \leftarrow o]; \dots; c_n[x \leftarrow o] & \\ o.m(v_1, \dots, v_n) \rightsquigarrow_{\epsilon} o' \implies C_p[o.m(v_1, \dots, v_n)] \rightsquigarrow C_p[o'] & \text{(external)} \end{array}$$

Fig. 3. Syntax, contexts and reduction rules of the data exploration calculus

3.1 Language syntax

The calculus combines object-oriented features such as member access with functional features including lambda functions. The syntax is defined in Figure 3. Object values o are defined by external libraries that are used in conjunction with the core calculus.

A program p in the data exploration calculus consists of a sequence of commands c . A command can be either a let binding or a term. Let bindings define variables x that can be used in subsequent commands. The calculus supports lambda functions, but they can only appear as arguments in method calls. A term t can be a value, variable or a member access, while an expression e , which can appear as an argument in member access, can be a lambda function or a term.¹

3.2 Operational semantics

The data exploration calculus is a call-by-value language. We model the evaluation as a small-step reduction \rightsquigarrow . Fully evaluating a program results in an irreducible sequence of objects $o_1; \dots; o_n$ (one object for each command, including let bindings) which can be displayed as intermediate results of the data analysis. The operational semantics is parameterized by a relation $\rightsquigarrow_{\epsilon}$ that models the functionality of the external libraries used with the calculus and defines reduction rules for member accesses. The relation has the following form:

$$o_1.m(v_1, \dots, v_n) \rightsquigarrow_{\epsilon} o_2$$

The operation is invoked on an object and takes values (objects or function values) as arguments. It always results in an opaque object. Figure 3 defines the reduction rules in terms of $\rightsquigarrow_{\epsilon}$ and evaluation contexts; C_e specifies left-to-right evaluation of arguments of a method call, C_c specifies evaluation of a command and C_p defines top-to-bottom evaluation of a program. The rule (external) applies reductions provided by an external library in a call-by-value order and (let) substitutes a value of an evaluated variable in all subsequent commands. Note that (let) leaves the value of the variable in the resulting list of commands.

¹Similar but weaker restrictions on the use of lambda functions exist in other languages. For example, in order to guide type inference, lambda functions in C# can appear as either method arguments or assigned to an explicitly typed variable, but they cannot be assigned to an implicitly typed variable.

3.3 Properties

The data exploration calculus is a very simple language and it has a number of desirable properties. However, some of those require that the relation $\rightsquigarrow_\epsilon$, which defines evaluation for external libraries, satisfies a number of conditions. Our main motivation is to capture properties that allow us to prove that our method of evaluating live previews, discussed in Section 6.1, is correct.

Definition 1 (Further reductions). We define two additional reduction relations:

- We write \rightsquigarrow^* for the reflexive, transitive closure of \rightsquigarrow
 - We write $\rightsquigarrow_{\text{let}}$ for a call-by-name let binding elimination
- $$c_1; \dots; c_{k-1}; \text{let } x = t; c_{k+1}; \dots; c_n \rightsquigarrow_{\text{let}} c_1; \dots; c_{k-1}; t; c_{k+1}[x \leftarrow t]; \dots; c_n[x \leftarrow t]$$

We say that two expressions e and e' are observationally equivalent if, for any context C , the expressions $C[e]$ and $C[e']$ reduce to the same value. Lambda functions $\lambda x \rightarrow 2$ and $\lambda x \rightarrow 1+1$ are not equal, but they are observationally equivalent. We require that external libraries satisfy two conditions. First, when a method is called with observationally equivalent values as arguments, it should return the same value. Second, the evaluation of $o.m(v_1, \dots, v_n)$ should be defined for all o, i and v_i . The external library will typically satisfy this by defining an error object and reducing all invalid calls to the error object, following the standard method of Milner [1978].

Definition 2 (External library). An external library consists of a set of objects O and a reduction relation $\rightsquigarrow_\epsilon$ that satisfies the following two properties:

- *Completeness*. For all o, m, i and all v_1, \dots, v_i , there exists o' such that $o.m(v_1, \dots, v_i) \rightsquigarrow_\epsilon o'$.
- *Compositionality*. For observationally equivalent arguments, the reduction returns the same object, i.e. given e_0, e_1, \dots, e_n and e'_0, e'_1, \dots, e'_n and m such that $e_0.m(e_1, \dots, e_n) \rightsquigarrow^* o$ and $e'_0.m(e'_1, \dots, e'_n) \rightsquigarrow^* o'$ then if for any contexts C_0, C_1, \dots, C_n it holds that if $C_i[e_i] \rightsquigarrow^* o_i$ and $C_i[e'_i] \rightsquigarrow^* o'_i$ for some o_i then $o = o'$.

The compositionality condition is essential for proving the correctness of our live preview mechanism. Completeness allows us to prove normalization, i.e. all programs reduce to a value – although the resulting value may be an error value provided by the external library.

Theorem 1 (Normalization). *For all p , there exists n and o_1, \dots, o_n such that $p \rightsquigarrow^* o_1; \dots; o_n$.*

PROOF. A program that is not a sequence of values can be reduced and reduction decreases the size of the program. See Appendix A.1 for more detail. \square

Although the reduction rules (let) and (external) of the data exploration calculus define an evaluation in a call-by-value order, eliminating let bindings in a call-by-name way using the $\rightsquigarrow_{\text{let}}$ reduction does not affect the result. This simplifies our later proof of live preview correctness in Section 6.1.

Lemma 2 (Let elimination for a program). *Given any program p such that $p \rightsquigarrow^* o_1; \dots; o_n$ for some n and o_1, \dots, o_n then if $p \rightsquigarrow_{\text{let}} p'$ for some p' then also $p' \rightsquigarrow^* o_1; \dots; o_n$.*

PROOF. The elimination of let binding transforms a program $c_1; \dots; c_{k-1}; \text{let } x = t; c_{k+1}; \dots; c_n$ to a program $c_1; \dots; c_{k-1}; t; c_{k+1}[x \leftarrow t]; \dots; c_n[x \leftarrow t]$. The reduction steps for the new program can be constructed using the steps of $p \rightsquigarrow^* o_1; \dots; o_n$. The new command t reduces to an object o using the same steps as the original term t in $\text{let } x = t$ but with context $C_c = -$ rather than $C_c = \text{let } x = -$; the terms t introduced by substitution also reduce using the same steps as before, but using contexts in which the variable x originally appeared. \square

4 FORMALISING LIVE CODING ENVIRONMENT

A naive way of providing live previews during code editing would be to re-evaluate the code after each change. This would be wasteful – for example, when writing code, most changes are additive and the preview can be updated by evaluating just the newly added code. In this section, we develop foundations for an efficient way of displaying live previews for the data exploration calculus.

4.1 Maintaining dependency graph

The key idea behind our method is to maintain a dependency graph [Kuck et al. 1981] with nodes representing individual operations of the computation that can be (partially) evaluated to obtain a preview. Each time the program text is modified, we parse it afresh (using an error-recovering parser) and bind the abstract syntax tree to the dependency graph. When binding a new expression to the graph, we reuse previously created nodes as long as they have the same dependencies. For expressions that have a new structure, we create new nodes.

The nodes of the graph serve as unique keys into a lookup table with previously evaluated parts of the program. When a preview is requested for an expression, we use the graph node bound to the expression to find a preview. If a preview has not been evaluated, we force the evaluation of all dependencies in the graph and then evaluate the operation represented by the current node.

4.1.1 Elements of the graph. The nodes of the graph represent individual operations of the computation. In our design, the nodes are used as cache keys, so we attach a unique symbol s to some of the nodes. That way, we can create two unique nodes representing, for example, access to a member named `take` which differ in their dependencies.

Furthermore, the graph edges are labeled with labels indicating the kind of dependency. For a method call, the labels are “first argument”, “second argument” and so on. Writing s for symbols and i for integers, nodes (vertices) v and edge labels l are defined as:

$$\begin{aligned} v &= \text{val}(o) \mid \text{var}(x) \mid \text{mem}(m, s) \mid \text{fun}(x, s) & (\text{Vertices}) \\ l &= \text{body} \mid \text{arg}(i) & (\text{Edge labels}) \end{aligned}$$

The `val` node represents a primitive value and contains the object itself. Two occurrences of `10` in the source code will be represented by the same node. Member access `mem` contains the member name, together with a unique symbol s – two member access nodes with different dependencies will contain a different symbol. Dependencies of member access are labeled with `arg` indicating the index of the argument (the instance has index 0 and arguments start with 1). Finally, nodes `fun` and `var` represent function values and variables bound by λ abstraction. For simplicity, we use variable names rather than de Bruijn indices and so renaming a bound variable forces recomputation.

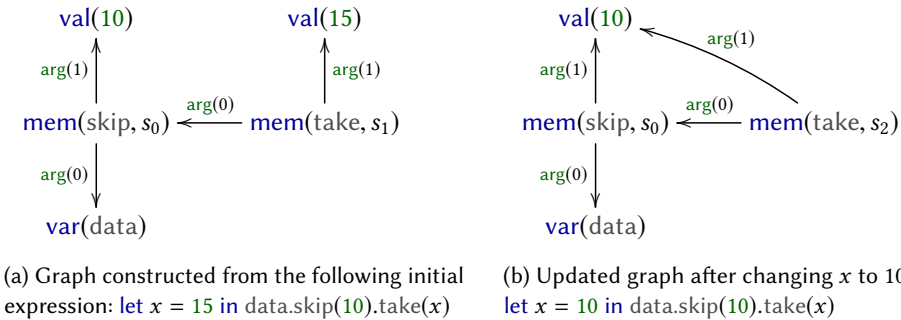


Fig. 4. Dependency graphs formed by two steps of the live programming process.

4.1.2 Example graph. Figure 4 illustrates how we construct and update the dependency graph. Node representing `take(x)` depends on the argument – the number 15 – and the instance, which is a node representing `skip(10)`. This, in turn, depends on the instance data and the number 10. Note that variables bound via `let` binding such as `x` do not appear as `var` nodes. The node using it depends directly on the node representing the expression that is assigned to `x`.

After changing the value of `x`, we create a new graph. The dependencies of the node `mem(skip, s0)` are unchanged and so the symbol `s0` attached to the node remains the same and previously computed previews can be reused. This part of the program is not recomputed. The `arg(1)` dependency of the `take` call changed and so we create a new node `mem(skip, s2)` with a fresh symbol `s2`. The preview for this node is then computed as needed using the already known values of its dependencies.

4.1.3 Reusing graph nodes. The binding process takes an expression and constructs a dependency graph reusing existing nodes when possible. For this, we use a lookup table of member access and function value nodes. The key in the lookup table is formed by a node kind (for disambiguation) together with a list of dependencies. A node kind is a member access containing the member name or a function containing the bound variable name; a lookup table Δ then maps a node kind with a list of dependencies to a cached node:

$$\begin{aligned} k &= \text{fun}(x) \mid \text{mem}(m) && (\text{Node kinds}) \\ \Delta(k, [(v_1, l_1), \dots, (v_n, l_n)]) && (\text{Lookup for a node}) \end{aligned}$$

The example on the second line looks for a node of a kind k that has dependencies v_1, \dots, v_n labeled with labels l_1, \dots, l_n . We write $\Delta(k, l) \downarrow$ when a value for a given key is not defined. For example, when creating the graph in Figure 4b, we perform the following lookup for the `skip` member access:

$$\Delta(\text{mem}(\text{skip}), [(\text{var}(\text{data}), \text{arg}(0)), (\text{val}(10), \text{arg}(1))])$$

The lookup returns the node `mem(skip, s0)` known from the previous step. We then perform the following lookup for the `take` member access:

$$\Delta(\text{mem}(\text{take}), [(\text{mem}(\text{skip}, s_0), \text{arg}(0)), (\text{val}(10), \text{arg}(1))])$$

In the previous graph, the argument of `take` was 15 rather than 10 and so this lookup fails. We then construct a new node `mem(take, s2)` using a fresh symbol `s2`.

4.2 Binding expressions to a graph

After parsing updated code, we update the dependency graph and link each node of the abstract syntax tree to a node of the dependency graph. This process is called binding and is defined by the `bind-expr` function (Figure 5) and `bind-prog` function (Figure 6). Both functions are annotated with a lookup table Δ discussed in Section 4.1 and a variable context Γ . The variable context is a map from variable names to dependency graph nodes and is used for variables bound using `let` binding.

When applied on an expression e , binding `bind-expr Γ, Δ (e)` returns a node v corresponding to the expression e paired with a dependency graph (V, E) . In the graph, V is a set of nodes v and E is a set of labeled edges (v_1, v_2, l) . We attach the label directly to the edge rather than keeping a separate colouring function as this makes the formalisation simpler. The `bind-prog Γ, Δ` function works similarly, but takes a sequence of commands and returns a sequence of nodes.

4.2.1 Binding expressions. When binding member access, we use `bind-expr` recursively to get a node and a dependency graph for each sub-expression. The nodes representing sub-expressions are then used as dependencies for lookup into Δ , together with their labels. If a node already exists in Δ it is reused (1). Alternatively, we create a new node containing a fresh symbol (2).

$$\begin{aligned}
\text{bind-expr}_{\Gamma, \Delta}(e_0.m(e_1, \dots, e_n)) &= v, (\{v\} \cup V_0 \cup \dots \cup V_n, E \cup E_0 \cup \dots \cup E_n) \\
&\quad \text{when } v_i, (V_i, E_i) = \text{bind-expr}_{\Gamma, \Delta}(e_i) \text{ and } v = \Delta(\text{mem}(m), [(v_0, \text{arg}(0)), \dots, (v_n, \text{arg}(n))]) \\
&\quad \text{let } E = \{(v, v_0, \text{arg}(0)), \dots, (v, v_n, \text{arg}(n))\} \tag{1} \\
\text{bind-expr}_{\Gamma, \Delta}(e_0.m(e_1, \dots, e_n)) &= v, (\{v\} \cup V_0 \cup \dots \cup V_n, E \cup E_0 \cup \dots \cup E_n) \\
&\quad \text{when } v_i, (V_i, E_i) = \text{bind-expr}_{\Gamma, \Delta}(e_i) \text{ and } \Delta(\text{mem}(m), [(v_0, \text{arg}(0)), \dots, (v_n, \text{arg}(n))]) \downarrow \\
&\quad \text{let } v = \text{mem}(m, s), s \text{ fresh and } E = \{(v, v_0, \text{arg}(0)), \dots, (v, v_n, \text{arg}(n))\} \tag{2} \\
\text{bind-expr}_{\Gamma, \Delta}(\lambda x \rightarrow e) &= v, (\{v\} \cup V, \{e\} \cup E) \\
&\quad \text{when } \Gamma_1 = \Gamma \cup \{x, \text{var}(x)\} \text{ and } v_0, (V, E) = \text{bind-expr}_{\Gamma_1, \Delta}(e) \text{ and } v = \Delta(\text{fun}(x), [(v_0, \text{body})]) \\
&\quad \text{let } e = (v, v_0, \text{body}) \tag{3} \\
\text{bind-expr}_{\Gamma, \Delta}(\lambda x \rightarrow e) &= v, (\{v\} \cup V, \{e\} \cup E) \\
&\quad \text{when } \Gamma_1 = \Gamma \cup \{x, \text{var}(x)\} \text{ and } v_0, (V, E) = \text{bind-expr}_{\Gamma_1, \Delta}(e) \text{ and } \Delta(\text{fun}(x), [(v_0, \text{body})]) \downarrow \\
&\quad \text{let } v = \text{fun}(x, s), s \text{ fresh and } e = (v, v_0, \text{body}) \tag{4} \\
\text{bind-expr}_{\Gamma, \Delta}(o) &= \text{val}(o), (\{\text{val}(o)\}, \emptyset) \tag{5} \\
\text{bind-expr}_{\Gamma, \Delta}(x) &= v, (\{v\}, \emptyset) \text{ when } v = \Gamma(x) \tag{6}
\end{aligned}$$

Fig. 5. Binding rules that define a construction of a dependency graph for an expression.

$$\begin{aligned}
\text{bind-prog}_{\Gamma, \Delta}(\text{let } x = e; c_2; \dots; c_n) &= v_1; \dots; v_n, (\{v_1\} \cup V \cup V_1, E \cup E_1) \\
&\quad \text{let } v_1, (V_1, E_1) = \text{bind-expr}_{\Gamma, \Delta}(e) \text{ and } \Gamma_1 = \Gamma \cup \{(x, v_1)\} \\
&\quad \text{and } v_2; \dots; v_n, (V, E) = \text{bind-prog}_{\Gamma_1, \Delta}(c_2; \dots; c_n) \tag{7} \\
\text{bind-prog}_{\Gamma, \Delta}(e; c_2; \dots; c_n) &= v_1; \dots; v_n, (\{v_1\} \cup V \cup V_1, E \cup E_1) \\
&\quad \text{let } v_1, (V_1, E_1) = \text{bind-expr}_{\Gamma, \Delta}(e) \text{ and } v_2; \dots; v_n, (V, E) = \text{bind-prog}_{\Gamma_1, \Delta}(c_2; \dots; c_n) \tag{8} \\
\text{bind-prog}_{\Gamma, \Delta}([]) &= [], (\emptyset, \emptyset) \tag{9}
\end{aligned}$$

Fig. 6. Binding rules that define a construction of a dependency graph for a program.

If a lambda function uses its argument, we will not be able to evaluate its body. In this case, the graph node bound to a function will depend on a synthetic node $\text{var}(x)$ that represents the variable with unknown value. When binding a function, we create the synthetic variable and add it to the variable context Γ_1 before binding the body. As with member access, the node representing a function may (3) or may not (4) be already present in the lookup table.

4.2.2 Binding programs. When binding a program, we bind the first command and then recursively process remaining commands until we reach an empty list of commands (9). For **let** binding (7), we bind the expression e assigned to the variable to obtain a graph node v_1 . We then bind the remaining commands using a variable context Γ_1 that maps the value of the variable to the graph node v_1 . The variable context is used when binding a variable (6) and so all variables declared using **let** binding will be bound to a graph node representing the value assigned to the variable. When the command is just an expression (8), we bind the expression using **bind-expr**.

$\text{update}_{V,E}(\Delta_{i-1}) = \Delta_i$ such that:

$$\begin{aligned} \Delta_i(\text{mem}(m), [(v_0, \text{arg}(0)), \dots, (v_n, \text{arg}(n))]) &= \text{mem}(m, s) \\ &\text{when } \text{mem}(m, s) \in V \text{ and } (\text{mem}(m, s), v_i, \text{arg}(i)) \in E \text{ for } i \in 0, \dots, n \\ \Delta_i(\text{fun}(x), [(v_1, \text{body})]) &= \text{fun}(x, s) \\ &\text{when } \text{fun}(x, s) \in V \text{ and } (\text{fun}(x, s), v_1, \text{body}) \in E \\ \Delta_i(v) &= \Delta_{i-1}(v) \quad (\text{otherwise}) \end{aligned}$$

Fig. 7. Updating the node cache after binding a new graph

4.3 Edit and rebind loop

The binding process formalised in Section 4.2 specifies how to update the dependency graph after updated program text is parsed. During live coding, this is done repeatedly as the programmer edits code. Throughout the process, we maintain a series of lookup table states $\Delta_0, \Delta_1, \Delta_2, \dots$. The initial lookup table is empty, i.e. $\Delta_0 = \emptyset$. At a step i , we parse a program p_i (consisting of several commands) and obtain a new dependency graph using the previous Δ . The result is a sequence of nodes corresponding to commands of the program and a graph (V, E) :

$$v_1; \dots; v_n, (V, E) = \text{bind-prog}_{\emptyset, \Delta_{i-1}}(p_i)$$

The new state of the cache is computed by calling the $\text{update}_{V,E}(\Delta_{i-1})$ function defined in Figure 7. The function adds newly created nodes from the graph (V, E) to the previous cache Δ_{i-1} and returns a new cache Δ_i . We only cache nodes for function and member accesses – nodes for variables and primitive values will remain the same thanks to the way they are constructed.

5 COMPUTING LIVE PREVIEWS

The binding process described in the previous section constructs a dependency graph after code changes. The nodes in the dependency graph correspond to individual operations that will be performed when running the program. When evaluating a preview, we attach (partial) results to nodes of the graph. Since the binding process reuses nodes when their dependencies do not change, previews for expressions (or sub-expressions) of a program can be reused when updating a preview.

In this section, we describe how previews are evaluated. The evaluation is done over the dependency graph, rather than over the structure of the program expressions as in the operational semantics given in Section 3.2. We analyse the preview evaluation formally in Section 6 and show that the resulting previews are the same as the result we would get by directly evaluating the code and, also, that no recomputation occurs when code is edited in certain ways.

5.1 Previews and delayed previews

Programs in the data exploration calculus consist of sequence of commands. Those can always be evaluated to a value with a preview that can be displayed to the user. However, we also support previews for all sub-expressions. This can be problematic if the current sub-expression is inside the body of a function. For example:

```
let top = movies.take(10).map( $\lambda x \rightarrow x.\text{getReleased}().\text{format}(\text{"dd-mm-yyyy"})$ )
```

Here, we can directly evaluate sub-expressions `movies` and `movies.take(10)`, but not `x.getReleased()` and `x.getReleased().format("dd-mm-yyyy")` because they contain a free variable `x`. Our preview evaluation algorithm addresses this by producing two kinds of previews. A *fully evaluated preview* is just a value, while a *delayed preview* is a partially evaluated expression with free variables:

$$\begin{aligned}
p &= o \mid \lambda x \rightarrow e && \text{(Fully evaluated previews)} \\
d &= p \mid \llbracket e \rrbracket_\Gamma && \text{(Evaluated and delayed previews)}
\end{aligned}$$

A fully evaluated preview p can be either a primitive object or a function value with no free variables. A possibly delayed preview d can be either an evaluated preview p or an expression e that requires variables Γ . For simplicity, we use an untyped language and so Γ is just a list of variables x_1, \dots, x_n .

A delayed preview is not necessarily the body of a lambda function as it appears in the source code. We partially evaluate sub-expressions of the body that do not have free variables or that have free variables bound by an earlier let binding.

Our implementation does not currently display delayed previews to the user, but there is a number of possible approaches for doing that. Most interestingly, since lambda functions always appear as arguments of a member access, we could force the evaluation of the surrounding expression, capture (a number of) values passed as inputs to the lambda function and display a preview based on those. Finally, in Section 7.2, we also consider a more speculative design for an abstraction mechanism that supports live previews that could, in some cases, replace lambda function.

5.2 Evaluation of previews

The evaluation of previews is defined in Figure 8. Given a dependency graph (V, E) , we define a relation $v \Downarrow d$ that evaluates a sub-expression corresponding to the node v to a possibly delayed preview d . The nodes V and edges E of the dependency graph are parameters of the \Downarrow relation, but they do not change during the evaluation and so we do not explicitly write them.

The auxiliary relation $v \Downarrow_{\text{lift}} d$ always evaluates to a delayed preview. If the ordinary evaluation returns a delayed preview, so does the auxiliary relation (lift-expr). If the ordinary evaluation returns a value, the value is wrapped into a delayed preview requiring no variables (lift-prev). A node representing a value is evaluated to a value (val) and a node representing an unbound variable is reduced to a delayed preview that requires the variable and returns its value (var).

$$\begin{array}{l}
\text{(lift-expr)} \quad \frac{v \Downarrow \llbracket e \rrbracket_\Gamma}{v \Downarrow_{\text{lift}} \llbracket e \rrbracket_\Gamma} \qquad \text{(fun-val)} \quad \frac{(\text{fun}(x, s), v, \text{body}) \in E \quad v \Downarrow p}{\text{fun}(x, s) \Downarrow \lambda x \rightarrow p} \\
\text{(lift-prev)} \quad \frac{v \Downarrow p}{v \Downarrow_{\text{lift}} \llbracket p \rrbracket_\emptyset} \qquad \text{(fun-bind)} \quad \frac{(\text{fun}(x, s), v, \text{body}) \in E \quad v \Downarrow \llbracket e \rrbracket_x}{\text{fun}(x, s) \Downarrow \lambda x \rightarrow e} \\
\text{(val)} \quad \frac{}{\text{val}(o) \Downarrow o} \qquad \text{(fun-expr)} \quad \frac{(\text{fun}(x, s), v, \text{body}) \in E \quad v \Downarrow \llbracket e \rrbracket_{x, \Gamma}}{\text{fun}(x, s) \Downarrow \llbracket \lambda x \rightarrow e \rrbracket_\Gamma} \\
\text{(var)} \quad \frac{}{\text{var}(x) \Downarrow \llbracket x \rrbracket_x} \\
\text{(mem-val)} \quad \frac{\forall i \in \{0 \dots k\}. (\text{mem}(m, s), v_i, \text{arg}(i)) \in E \quad v_i \Downarrow p_i \quad p_0.m(p_1, \dots, p_k) \rightsquigarrow_\epsilon p}{\text{mem}(m, s) \Downarrow p} \\
\text{(mem-expr)} \quad \frac{\forall i \in \{0 \dots k\}. (\text{mem}(m, s), v_i, \text{arg}(i)) \in E \quad \exists j \in \{0 \dots k\}. v_j \not\Downarrow p_j \quad v_i \Downarrow_{\text{lift}} \llbracket e_i \rrbracket_{\Gamma_i}}{\text{mem}(m, s) \Downarrow \llbracket e_0.m(e_1, \dots, e_k) \rrbracket_{\Gamma_0, \dots, \Gamma_k}}
\end{array}$$

Fig. 8. Rules that define evaluation of previews over a dependency graph for a program

For member access, we distinguish two cases. If all arguments evaluate to values (member-val), then we use the evaluation relation defined by external libraries \rightsquigarrow_e to immediately evaluate the member access and produce a value. If some of the arguments are delayed (member-expr), because the member access is inside the body of a lambda function, we produce a delayed member access expression that requires the union of the variables required by the individual arguments.

The evaluation of function values is similar, but requires three cases. If the body can be reduced to a value with no unbound variables (fun-val), we return a lambda function that returns the value. If the body requires only the bound variable (fun-bind), we return a lambda function with the delayed preview as the body. If the body requires further variables, the result is a delayed preview.

5.3 Caching of evaluated previews

For simplicity, the relation \Downarrow in Figure 8 does not specify how previews are cached and linked to graph nodes. In practice, this is done by maintaining a lookup table from graph nodes v to (possibly delayed) previews p . Whenever \Downarrow is used to obtain a preview for a graph node, we first attempt to find an already evaluated preview using the lookup table. If the preview has not been previously evaluated, we evaluate it and add it to the lookup table.

The cached evaluated previews can be reused in two ways. First, multiple nodes can depend on one sub-graph in a single dependency graph (if the same sub-expression appears twice in the program). Second, the keys of the lookup table are graph nodes and nodes are reused when a new dependency graph is constructed after the user edits the source code.

5.4 Theories of delayed previews

The operational semantics presented in this paper serves two purposes. It gives a simple guide for implementing text-based live coding environments for data science and we use it to prove that our optimized way of producing live previews is correct. However, some aspects of our mechanism are related to important work in semantics of programming languages and deserve to be mentioned.

The construction of delayed previews is related to meta-programming. Assuming we have delayed previews $\llbracket e_0 \rrbracket_x$ and $\llbracket e_1 \rrbracket_y$ and we invoke a member m on e_0 using e_1 as an argument. To do this, we construct a new delayed preview $\llbracket e_0.m(e_1) \rrbracket_{x,y}$. This operation is akin to expression splicing from meta-programming [Syme 2006; Taha and Sheard 2000].

The semantics of delayed previews can be more formally captured by Contextual Modal Type Theory (CMTT) [Nanevski et al. 2008] and comonads [Gabbay and Nanevski 2013]. In CMTT, $[\Psi]A$ denotes that a proposition A is valid in context Ψ , which is similar to our delayed previews written as $\llbracket A \rrbracket_\Psi$. CMTT defines rules for composing context-dependent propositions that would allow us to express the splicing operation used in (mem-expr). In categorical terms, the context-dependent proposition can be modeled as an indexed comonad [Gaborardi et al. 2016; Mycroft et al. 2016]. The evaluation of a preview with no context dependencies (built implicitly into our evaluation rules) corresponds to the counit operation of a comonad and would be explicitly written as $\llbracket A \rrbracket_\emptyset \rightarrow A$.

6 EVALUATING LIVE CODING ENVIRONMENT

Evaluating previews using a dependency graph makes it possible to cache partial results. The mechanism satisfies two properties. First, if we evaluate a preview using dependency graph with caching, it is the same as the value we would obtain by evaluating the expression directly. Second, the evaluation of previews using dependency graphs reuses – in some cases – previously evaluated partial results. In other words, we show that the mechanism is correct and implements an effective optimization. We discuss correctness first and then evaluate effectiveness, both theoretically (Section 6.2) and empirically (Section 6.3). Finally, we describe a case study where we used the methods outlined in the paper to develop a larger online service for data exploration (Section 6.4).

6.1 Correctness of previews

To show that the previews are correct, we prove two properties. Correctness (Theorem 6) guarantees that, no matter how a dependency graph is constructed, when we use it to evaluate previews for a program, the previews are the same as the values we would obtain by evaluating the program commands directly. Determinacy (Theorem 7) guarantees that if we cache a preview for a graph node and update the graph, the preview we would evaluate using the updated graph would be the same as the cached preview.

To simplify the proofs, we consider programs without let bindings. This is possible, because eliminating let bindings does not change the result of evaluation, as shown earlier in Lemma 2, and it also does not change the constructed dependency graph as shown below in Lemma 3.

Lemma 3 (Let elimination for a dependency graph). *Given programs p_1, p_2 such that $p_1 \rightsquigarrow_{\text{let}} p_2$ and a lookup table Δ_0 then if $v_1; \dots; v_n, (V, E) = \text{bind-prog}_{\emptyset, \Delta_0}(p_1)$ and $v'_1; \dots; v'_n, (V', E') = \text{bind-prog}_{\emptyset, \Delta_1}(p_2)$ such that $\Delta_1 = \text{update}_{V, E}(\Delta_0)$ then for all i , $v_i = v'_i$ and also $(V, E) = (V', E')$.*

PROOF. Assume $p_1 = c_1; \dots; c_{k-1}; \text{let } x = e; c_{k+1}; \dots; c_n$ and the let binding is eliminated resulting in $p_2 = c_1; \dots; c_{k-1}; e; c_{k+1}[x \leftarrow e]; \dots; c_n[x \leftarrow e]$. When binding p_1 , the case $\text{bind-prog}_{\Gamma, \Delta}(\text{let } x = e)$ is handled using (7) and the node resulting from binding e is added to the graph V, E . It is then referenced each time x appears in subsequent commands $c_{k+1}; \dots; c_n$. When binding p_2 , the node resulting from binding e is a primitive value or a node already present in Δ_1 (added by $\text{update}_{V, E}$) and is reused each time $\text{bind-expr}_{\Gamma, \Delta_1}(e)$ is called. \square

The Lemma 3 provides a way of removing let bindings from a program, such that the resulting dependency graph remains the same. Here, we bind the original program first, which adds the node for e to Δ . In our implementation, this is not needed because Δ is updated while the graph is being constructed using bind-expr . To keep the formalisation simpler, we separate the process of building the dependency graph and updating Δ and thus Lemma 3 requires an extra binding step.

Now, we can show that, given a let-free expression, the preview obtained using a correctly constructed dependency graph is the same as the one we would obtain by directly evaluating the expression. This requires a simple auxiliary lemma.

Lemma 4 (Lookup inversion). *Given Δ obtained using update as defined in Figure 7 then:*

- If $v = \Delta(\text{fun}(x), [(v_0, l_0)])$ then $v = \text{fun}(x, s)$ for some s .
- If $v = \Delta(\text{mem}(m), [(v_0, l_0), \dots, (v_n, l_n)])$ then $v = \text{mem}(m, s)$ for some s .

PROOF. By construction of Δ in Figure 7. \square

Theorem 5 (Term preview correctness). *Given a term t that has no free variables, together with a lookup table Δ obtained from any sequence of programs using bind-prog (Figure 6) and update (Figure 7), then let $v, (V, E) = \text{bind-expr}_{\emptyset, \Delta}(t)$.*

If $v \Downarrow p$ over a graph (V, E) then $p = o$ for some value o and $t \rightsquigarrow^ o$.*

PROOF. First note that, when combining recursively constructed sub-graphs, the bind-expr function adds new nodes and edges leading from those new nodes. Therefore, an evaluation using \Downarrow over a sub-graph will also be valid over the new graph – the newly added nodes and edges do not introduce non-determinism to the rules given in Figure 8.

We prove a more general property showing that for any e , its binding $v, (V, E) = \text{bind-expr}_{\emptyset, \Delta}(e)$ and any evaluation context C such that $C[e] \rightsquigarrow o$ for some o , one of the following holds:

- a. If $FV(e) = \emptyset$ then $v \Downarrow p$ for some p and $C[p] \rightsquigarrow o$
- b. If $FV(e) \neq \emptyset$ then $v \Downarrow \llbracket e_p \rrbracket_{FV(e)}$ for some e_p and $C[e_p] \rightsquigarrow o$

In the first case, p is a value, but it is not always the case that $e \rightsquigarrow^* p$, because p may be lambda function and preview evaluation may reduce sub-expression in the body of the function. Using a context C in which the value reduces to an object avoids this problem.

The proof of the theorem follows from the more general property. Using a context $C[-] = -$, the term t reduces $t \rightsquigarrow^* t' \rightsquigarrow_{\epsilon} o$ for some o and the preview p is a value o because $C[p] = p = o$. The proof is by induction over the binding process, which follows the structure of the expression. The full proof is shown in Appendix A.2. \square

The correctness theorem combines the previous two results.

Theorem 6 (Program preview correctness). *Consider a program $p = c_1; \dots; c_n$ that has no free variables, together with a lookup table Δ_0 obtained from any sequence of programs using bind-prog (Figure 6) and update (Figure 7). Assume a let-free program $p' = t_1; \dots; t_n$ such that $p \rightsquigarrow_{\text{let}}^* p'$.*

Let $v_1; \dots; v_n, (V, E) = \text{bind-prog}_{\emptyset, \Delta_0}(p)$ and define updated lookup table $\Delta_1 = \text{update}_{V, E}(\Delta_0)$ and let $v'_1; \dots; v'_n, (V', E') = \text{bind-prog}_{\emptyset, \Delta_1}(p)$.

If $v'_i \Downarrow p_i$ over a graph (V', E') then $p_i = o_i$ for some value o_i and $t_i \rightsquigarrow o_i$.

PROOF. Direct consequence of Lemma 3 and Theorem 5. \square

As discussed earlier, our implementation updates Δ during the recursive binding process and so a stronger version of the property holds – namely, previews calculated over a graph obtained directly for the original program p are the same as the values of the fully evaluated program. We note that this is the case, but do not show it formally to aid the clarity of our formalisation.

The second important property that guarantees that our mechanism always displays correct previews is determinacy. This makes it possible to cache the previews evaluated using \Downarrow using the corresponding graph node as a lookup key.

Theorem 7 (Program preview determinacy). *For some Δ and for any programs p, p' , assume that the first program is bound, i.e. $v_1; \dots; v_n, (V, E) = \text{bind-prog}_{\emptyset, \Delta}(p)$, the graph node cache is updated $\Delta' = \text{update}_{V, E}(\Delta)$ and the second program is bound, i.e. $v'_1; \dots; v'_m, (V', E') = \text{bind-prog}_{\emptyset, \Delta'}(p')$. Now, for any v , if $v \Downarrow p$ over (V, E) then also $v \Downarrow p$ over (V', E') .*

PROOF. By induction over \Downarrow over (V, E) , we show that the same evaluation rules also apply over (V', E') . This is the case, because graph nodes added to Δ' by $\text{update}_{V, E}$ are added as new nodes in $\text{bind-prog}_{\emptyset, \Delta'}$ and nodes and edges of (V, E) used during the evaluation are unaffected. \square

The mechanism used for caching previews, as discussed in Section 5.3, keeps a preview or a partial preview d in a lookup table indexed by nodes v . The Theorem 7 guarantees that this is a valid strategy. As we update dependency graph during code editing, previous nodes will continue representing the same sub-expressions.

6.2 Reuse of previews

In the motivating example given in Section 1, the data analyst first extracted a constant value into a let binding and then modified a parameter of the last method call in a call chain. We argued that a live coding environment should reuse partially evaluated previews for these two cases. We now show that this is, indeed, the case in our system.

In this section, we identify a number of code edit operations where the previously evaluated values for a sub-expression can be reused. The list of operations is shown in Figure 9. To express the operations we define an editing context K which is defined similarly to evaluation context C from Figure 3, but captures sub-expressions appearing anywhere in the program.

Edit contexts of expressions

$$\begin{aligned} K_e[-] &= K_e[-].m(e_1, \dots, e_n) \mid e.m(e_1, \dots, e_{l-1}, K_e[-], e_{l+1}, \dots, e_n) \mid - \\ K_c[-] &= \text{let } x = K_e[-] \mid K_e[-] \end{aligned}$$

Code edit operations preserving preview for a sub-expression

- (let-intro-var) $\bar{c}_1; \langle e \rangle; \bar{c}_2$ changes to $\bar{c}_1; \text{let } x = e; \langle x \rangle; \bar{c}_2$ where x is a fresh name.
- (let-intro-ins) $\bar{c}_1; \bar{c}_2; \langle K_c[e] \rangle; \bar{c}_3$ is changed to $\bar{c}_1; \text{let } x = e; \bar{c}_2; \langle K_c[x] \rangle; \bar{c}_3$ via a semantically non-equivalent expression $\bar{c}_1; \bar{c}_2; K_c[x]; \bar{c}_3$ where x is free.
- (let-intro-del) $\bar{c}_1; \bar{c}_2; \langle K_c[e] \rangle; \bar{c}_3$ is changed to $\bar{c}_1; \text{let } x = e; \bar{c}_2; \langle K_c[x] \rangle; \bar{c}_3$ via an expression $\bar{c}_1; \text{let } x = e; \bar{c}_2; K_c[e]; \bar{c}_3$ with unused variable x .
- (let-elim-del) $\bar{c}_1; \text{let } x = e; \bar{c}_2; \langle K_c[x] \rangle; \bar{c}_3$ is changed to $\bar{c}_1; \bar{c}_2; \langle K_c[e] \rangle; \bar{c}_3$ via a semantically non-equivalent expression $\bar{c}_1; \bar{c}_2; K_c[x]; \bar{c}_3$ where x is free.
- (let-elim-ins) $\bar{c}_1; \text{let } x = e; \bar{c}_2; \langle K_c[x] \rangle; \bar{c}_3$ is changed to $\bar{c}_1; \bar{c}_2; \langle K_c[e] \rangle; \bar{c}_3$ via an expression $\bar{c}_1; \text{let } x = e; \bar{c}_2; K_c[e]; \bar{c}_3$ with unused variable x .
- (edit-mem) $\bar{c}_1; K_c[\langle e_0 \rangle.m(\bar{e})]; \bar{c}_2$ is changed to $\bar{c}_1; K_c[\langle e_0 \rangle.m'(\bar{e}')]; \bar{c}_2$
- (edit-let) $\bar{c}_1; \text{let } x = e_1; \bar{c}_2; K_c[\langle e_2 \rangle]; \bar{c}_3$ is changed to $\bar{c}_1; \text{let } x = e'_1; \bar{c}_2; K_c[\langle e_2 \rangle]; \bar{c}_3$ when $x \notin FV(e_2)$.

Fig. 9. Code edit operations that preserve previously evaluated preview

We use the notation $\langle e \rangle$ to mark the part of the expressions before and after the edit that have the same value that will not need to be recomputed after the change and we write \bar{c} and \bar{e} for a list of commands and expressions, respectively. In some of the edit operations, we also specify an intermediate program that is, in some cases, semantically different and only has a partial preview. This illustrates a typical way of working with code in a text editor using cut and paste operations. For example, in (let-intro-ins), we assume that the analyst cuts a sub-expression e , replaces it with a variable x and then adds a let binding for a variable x and inserts the expression e from the clipboard. Edit with the same result is captured by (let-intro-del) but here the analyst first inserts the let binding and then replaces the sub-expression e with a variable x .

The of preview-preserving code edit operations includes several ways of introducing a let binding, several ways of eliminating a let binding and two edit operations where the data analyst modifies an unrelated part of the program. It is worth noting that the list is not exhaustive. Rather, it aims to illustrate some of the typical operations that the data analyst might perform when writing code.

Theorem 9 proves that the operations given in Figure 9 preserve the preview for a marked sub-expression. It relies on the following lemma that characterizes one of the common kinds of edits more generally. Given two versions of a program that both contain the same sub-expression e , if the let bindings that define the values of variables used in e do not change, then the graph node assigned to e will be the same when binding the original and the updated program.

Lemma 8 (Binding sub-expressions). *Assume we have programs $p_1 = c_1; \dots; c_k; K_c[e]; c_{k+1}; \dots; c_n$ and $p_2 = c'_1; \dots; c'_k; K'_c[e]; c'_{k+1}; \dots; c'_n$ and $I \subseteq \{1 \dots k\}$ such that $\forall i \in I. c_i = c'_i$ and for each $x \in \bigcup_{i \in I} FV(c_i) \cup FV(e)$ there exists $j \in I$ such that $c_j = \text{let } x = e$ for some e . Given any Δ , assume that the first program is bound, i.e. $v_1; \dots; v_n, (V, E) = \text{bind-prog}_{0, \Delta}(p_1)$, the cache is updated $\Delta' = \text{update}_{V, E}(\Delta)$ and the second program is bound, i.e. $v'_1; \dots; v'_n, (V', E') = \text{bind-prog}_{0, \Delta'}(p_2)$.*

Now, assume $v, G = \text{bind-expr}_{\Gamma, \Delta}(e)$ and $v', G' = \text{bind-expr}_{\Gamma', \Delta'}(e)$ are the recursive calls to bind e during the first and the second binding, respectively. Then, the graph nodes assigned to the sub-expression e are the same, i.e. $v = v'$.

PROOF. First, assuming that $\forall x \in FV(e). \Gamma(x) = \Gamma'(x)$, we show by induction over the binding process of e for the first program that the result is the same. In cases (1) and (3), the updated Δ' contains the required key and so the second binding proceeds using the same case. In cases (2) and (4), the second binding reuses the node created by the first binding using case (1) and (3), respectively. Cases (5) and (6) are the same.

Second, when binding let bindings in $c_1; \dots; c_k$, the initial $\Gamma = \emptyset$ during both bindings. Nodes added to Γ and Γ' for commands c_j such that $j \in I$ are the same and nodes added for remaining commands do not add any new nodes referenced from e and so $v = v'$ using the above. \square

Theorem 9 (Preview reuse). *Given the sequence of expressions as specified in Figure 9, if the expressions are bound in sequence and graph node cache updated as specified in Figure 7, then the graph nodes assigned to the specified sub-expressions are the same.*

PROOF. Cases (edit-let) and (edit-mem) are direct consequences of Lemma 8; for (let-intro-var), the node assigned to x is the node assigned to e which is the same as before the edit from Lemma thm:sub-expr. Cases (let-intro-ins) and (let-intro-del) are similar to (let-intro-var), but also require using induction over the binding of $K_c[e]$. Finally, cases (let-elim-ins) and (let-elim-del) are similar and also use Lemma 8 together with induction over the binding of $K_c[x]$. \square

6.3 Empirical evaluation of efficiency

The key performance claim about our method of providing instantaneous feedback is that it is more efficient than recomputing values for the whole program (or the current command) after every keystroke. In the previous section, we formally proved that this is true and gave examples of code edit operations that do not cause recomputation. In this section, we further support this claim with an empirical evaluation. The purpose of this section is not to precisely evaluate overheads of our implementation, but to compare how much recomputation different evaluation strategies perform.

For the purpose of the evaluation, we use a simple image manipulation library that provides operations for loading, greyscaling, blurring and combining images. We compare delays in updating the preview for three different evaluation strategies, while performing the same sequence of code edit operations. Using image processing as an example gives us a way to visualize the reuse of previously computed values. As in a typical data exploration scenario, the individual operations are relatively expensive compared to the overheads of building the dependency graph.

6.3.1 Code edit operations and methodology. To avoid clutter when visualizing the performance, we update the preview after complete tokens are added (such as identifier, number, dot or left parenthesis) rather than after individual keystrokes. Figure 10 shows the sequence of code edit operations that we use to measure the delays in updating a live preview. We first enter an expression to load, greyscale and blur an image (1) then introduce let binding (2) and add more operations (3). Finally, we extract one of the parameters into a variable (4). Most of the operations are simply adding code, but there are two cases where we modify existing code and change value of a parameter for blur and combine immediately after (1) and (3), respectively.

We implement the algorithm described in Sections 4 and 5 in a simple web-based environment that allows the user to modify code and explicitly trigger recomputation. It then measures time needed to recompute values for the whole program and displays the resulting image. If the parsing fails, e.g. because of an unclosed parenthesis, previews are not updated and we record only the time taken by parsing attempt. We compare the delays of three different evaluation strategies:

- (1) Enter the following code and then change parameter of blur from 4 to 8:

```
image.load("shadow.png").greyScale().blur(4)
```

- (2) Assign the result to a variable and start writing code for further operations:

```
let shadow = image.load("shadow.png").greyScale().blur(8)
shadow.combine
```

- (3) Finish code to combine two images and then change parameter of combine from 20 to 80:

```
let shadow = image.load("shadow.png").greyScale().blur(8)
shadow.combine(image.load("pope.png"), 20)
```

- (4) Extract the parameter of combine to a let bound variable:

```
let ratio = 80
let shadow = image.load("shadow.png").greyScale().blur(8)
shadow.combine(image.load("pope.png"), ratio)
```

Fig. 10. Code edit operations that are used in the experimental evaluation

- *Call-by-value*. Following the semantics in Section 3.2, all sub-expressions are fully evaluated before an expression is evaluated. In this case, work may be done even if the evaluation results in an error. For example, we parse `image.load("shadow.png").blur` as member access with no arguments. The evaluation loads the image, but then fails because `blur` requires one argument.
- *Lazy*. To address the wastefulness of call-by-value strategy discussed above, we simulate lazy evaluation by implementing a version of the image processing library where operations build a delayed computation and only evaluate it when rendering an image. Using this strategy, failing computations do not perform unnecessary work.
- *Live*. Finally, the live strategy uses the algorithm described in Sections 4 and 5. The cache is empty at the beginning of the experiment and we update it after each token is added. This is the only strategy that does not start afresh after reparsing code.

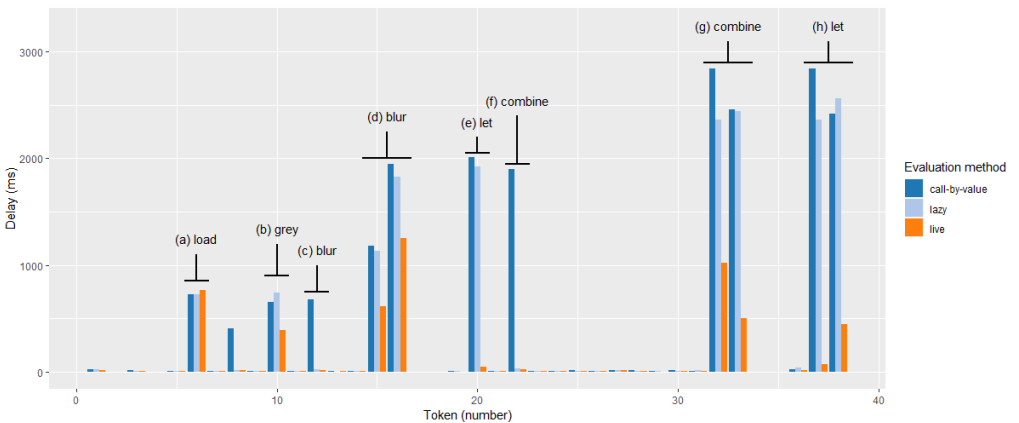


Fig. 11. Time required to recompute the results of a sample program after individual tokens of the program are added or modified for three different evaluations strategies.




Fig. 12. Histogram showing the distribution of delays incurred when updating previews. We show a histogram computed from all delays (left) and only from delays larger than 15ms (right).

6.3.2 Efficiency evaluation. The experimental environment is implemented in F# and compiled to JavaScript using the Fable compiler. We run the experiments in Firefox (version 64.0.2, 32bit) on Windows 10 (build 1809, 64bit) on a computer with Intel Core i7-7660U CPU and 16Gb RAM.

Figure 11 shows times needed to recompute previews after individual tokens of the program are added, deleted or modified, according to the script in Figure 10, resulting in 38 measurements for each of the 3 evaluation strategies. We mark a number of notable points in the chart:

- (a) Loading image for the first time incurs small extra overhead in the live strategy.
- (b) Greyscaling using the live strategy does not need to re-load the image.
- (c) Accessing the blur member without arguments causes delay for call-by-value strategy.
- (d) When varying the parameter of blur, the live strategy reuses previously greyscaled image.
- (e) Introducing let binding does not cause recomputation when using live strategy.
- (f) As in (c), accessing combine member without valid argument only affects call-by-value.
- (g) Varying parameter of combine in live strategy does not recompute blur and is much faster.
- (h) Introducing let binding, again, causes full recomputation for lazy and call-by-value.

A summarized view of the delays is provided in Figure 12, which shows histograms illustrating the distribution of delays for each of the three evaluation methods. A large proportion of delays is very small (less than 15ms) because the parser used in our experimental environment often fails (e.g. for unclosed parentheses). The histogram on the right summarizes only delays for edit operations where the delay for at least one of the strategies was over 15ms. A more robust parser would be able to recover and the delays for call-by-value and lazy strategies would be even more significant.

6.3.3 Summary of results. The purpose of our empirical is not to exactly assess the overhead of our implementation of the live evaluation strategy algorithm. Instead, our goal is to illustrate how often can previously evaluated results be reused and the impact this has when writing code. The experiment presented in this section is small-scale, but it is sufficient to illustrate the main point.

When recomputing results after every edit operation using the *call-by-value* strategy, the time needed to update the results grows continually. The *lazy* strategy removes the overhead for programs that fail to evaluate, but keeps the same trend. Using our *live* strategy, the computation can reuse values computed previously and, as a result, expensive operations (such as (d) and (g) in Figure 11) are significantly faster, because they do not need to recompute operations done previously when writing the code. As shown in Figure 12, there are almost no very expensive operations (taking over 1 second) in the *live* strategy in contrast to several in the other two strategies.

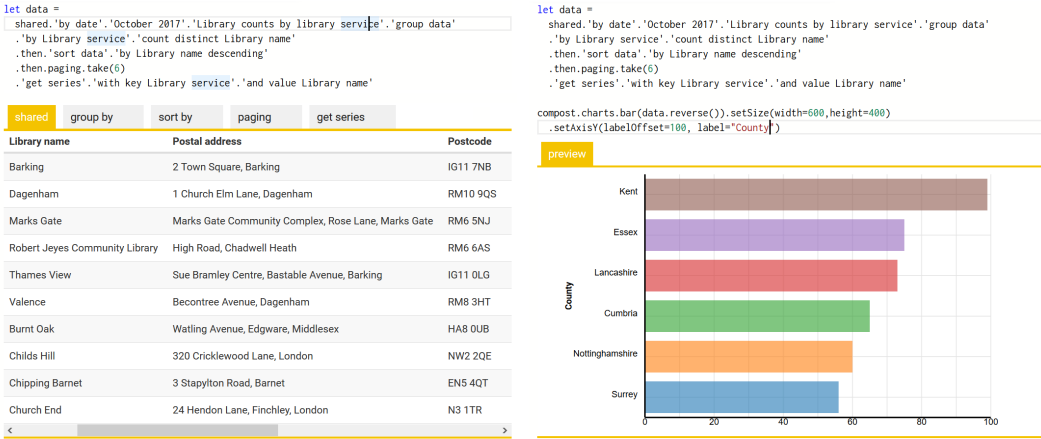


Fig. 13. Data analysis counting the number of libraries per county in the UK. We load and aggregate data (left) and then create a chart with a label (right). The example is adapted from <http://gallery.thegamma.net/73/>.

6.4 Transparent tools for data journalism

In Section 2.1, we motivated our work by considering how journalists explore open data. In addition to the theoretical and experimental work presented in this paper, we also implemented an online data exploration environment, equipped with live editor for The Gamma language that provides instantaneous feedback during coding. The environment uses the principles presented in this paper to build a more comprehensive system that allows users, such as journalists, to analyse, summarize and visualize open data. In this section, we briefly report on our experience with the system.

The environment is available at <http://gallery.thegamma.net>. Two screenshots shown in Figure 13 illustrate a number of interesting features:

- The left screenshot uses a type provider for data aggregation introduced by Petricek [2017b] and briefly discussed in Section 2.2. For the purpose of this paper, the type provider is an external library (with objects, members and reduction relation). Using the type provider relies on editor auto-complete which is based on type information. To support this, we implement type checking, as discussed in Section 7.1, reusing the mechanism of live preview evaluation.
- The live preview for code written using the data aggregation type provider (left) is shown in a panel below the current command. This shows the preview for a sub-expression based on the position of the cursor (i.e. raw data before grouping in the screenshot). In addition to the preview, our environment also generates tabs that let the user navigate to individual steps of the data transformation for code using the data aggregation type provider.
- The right screenshot uses a charting library, which is another external library available in the environment. Unlike for data aggregation (left), live previews for charts (right) are not handled in a special way and we show the result of evaluating the whole command. The screenshot shows a case where the user modifies parameters of the chart. Thanks to our live evaluation strategy, this is done efficiently without reevaluating the data transformation.

Conducting an extensive user study that would evaluate the usability of the system from a human-computer interaction perspective is left for future work. To give anecdotal evidence of the usability, the example shown in Figure 13 is an analysis developed by an attendee of our tutorial at Mozilla Festival 2017 who had no prior programming experience.




Fig. 14. Dependency graph for $o.m(\lambda x \rightarrow x)$ with a newly added **callsite** edges.

7 EXTENDING DATA EXPLORATION CALCULUS

The data exploration calculus presented in Section 3 allowed us to formalise the live preview mechanism, but it lacks many important features. We outline two extensions. First, we discuss how to add instantaneous type checking by using the dependency graph construction described in Section 4. Second, we consider an abstraction mechanism that preserves the live coding experience.

7.1 Type checking

Instantaneous live previews give analysts quick feedback when they write incorrect code, but having type information is still valuable. First, it can help give better error messages. Second, types can be used to provide auto-complete – when the user types ‘.’ we can offer available members without having to evaluate the wait until the value of the object is available.

7.1.1 Revised dependency graph. Type checking of small programs is typically fast enough that no caching is necessary. However, The Gamma supports *type providers* [Christiansen 2013; Syme et al. 2013], which can generate types based on an external file or a REST service call, e.g. [Petricek et al. 2016]. For this reason, type checking can be relatively time consuming and can benefit from the same caching facilities as those available for live previews.

Adding type checking requires revising the way we construct the dependency graph introduced in Section 4. Previously, a variable bound by a lambda function had no dependencies. However, the type of the variable depends on the context in which it appears. Given an expression $o.m(\lambda x \rightarrow x)$, we infer the type of x from the type of the first argument of the member m . A variable node for x thus needs to depend on the call site of m . We capture that by adding an edge **callsite**(m, i) from x to o which indicates that x is the input variable of a function passes as the i^{th} argument to the m member of the expression represented by the target node. We also add **callsite**(m, i) as an edge from the node of the function. Figure 14 shows the revised dependency graph for $o.m(\lambda x \rightarrow x)$. A revised version of the algorithm that constructs the dependency graph is included in Appendix B.

7.1.2 Type checking. The structure of typing rules is similar to the evaluation relation $v \Downarrow d$ defined earlier. Given a dependency graph (V, E) , we define typing judgements in the form $v \vdash \tau$. The type τ can be a primitive type, a function $\tau \rightarrow \tau$ or an object type $\{m_1 : \sigma_1, \dots, m_n : \sigma_n\}$ with member types $\sigma = (\tau_1, \dots, \tau_n) \rightarrow \tau$.

The typing rules for variables, functions and member access are shown in Figure 15. Remaining rules are available in Appendix B. When type checking a member access (**mem**), we find its dependencies v_i and check that the instance is an object with the required member m . The types of arguments of the member then need to match the types of the remaining (non-instance) nodes. Type checking a function (**fun**) and a variable (**var**) is similar. In both cases, we follow the **callsite** edge to find the member that accepts the function as an argument and use the type of the argument to check the type of the function or infer the type of the variable.

The results of type checking can be cached and reused in the same way as live previews, although we leave out the details. A property akin to correctness (Theorem 6) requires defining standard type checking over the structure of expressions, which we also omit for space reasons.

$$\begin{array}{c}
(\text{mem}) \frac{\forall i \in \{0 \dots k\}. (\text{mem}(m, s), v_i, \text{arg}(i)) \in E \quad v_0 \vdash \{.., m : (\tau_1, \dots, \tau_k) \rightarrow \tau, ..\} \quad v_i \vdash \tau_i}{\text{mem}(m, s) \vdash \tau} \\
\\
(\text{var}) \frac{(\text{var}(x, s), v, \text{callsite}(m, i)) \in E \quad v \vdash \{.., m : (\tau_1, \dots, \tau_k) \rightarrow \tau, ..\} \quad \tau_i = \tau' \rightarrow \tau''}{\text{var}(x, s) \vdash \tau'} \\
\\
(\text{fun}) \frac{\{(\text{fun}(x, s), v_b, \text{body}), (\text{var}(x, s), v_c, \text{callsite}(m, i))\} \subseteq E \quad v_c \vdash \{.., m : (\tau_1, \dots, \tau_k) \rightarrow \tau, ..\} \quad \tau_i = \tau' \rightarrow \tau'' \quad v_b \vdash \tau''}{\text{fun}(x, s) \vdash \tau' \rightarrow \tau''}
\end{array}$$

Fig. 15. Rules that define type checking of terms and expressions over a dependency graph (V, E)

7.2 Feedback-friendly abstraction

The sample data analysis by Financial Times in Section 2.1 illustrates why notebook users often avoid abstraction. Wrapping code in a function would make it impossible to split code into cells, and see results of intermediate steps. Instead, the sample analysis used a global variable at the top of the notebook with possible values in a comment.

Providing live previews inside ordinary functions is problematic, because we do not have readily available values for input parameters and our mechanism for lambda functions only provides delayed previews inside body of a function. We believe that extending the data exploration calculus with an abstraction mechanism that would support development with instantaneous feedback is an interesting design problem and we briefly outline a possible solution here.

Data scientists often write code interactively using a sample data set and, when it works well, wrap it into a function that they then call on other data. Similarly, spreadsheet users often write equation in the first row of a table and then use the “drag down” operation to apply it to other rows. One way of adding similar functionality to the data exploration calculus is to label a sequence of commands such that the sequence can be reused later with different inputs:

$$\begin{array}{lcl}
p & = & c_1; \dots; c_n \\
c & = & \text{let } x = t \mid t \mid \text{lbl}: p \mid \text{lbl}
\end{array}$$

As before, a program is a sequence of commands and a command can be either a let binding or a term. Additionally, a command can be a labelled sequence of commands or a reference to a label. When evaluating, the command *lbl* is replaced with the associated sequence of commands *p* before any other reduction occurs. Consequently, variables used in the labelled block are dynamically scoped and we can use let binding to redefine a value of a variable before invoking the block repeatedly. The correct use of dynamically scoped variables can be checked using coeffects [Petricek 2017a].

This minimalistic abstraction mechanism supports code reuse without affecting how live previews are computed. The commands in the labelled block still require all let bound variables to be defined before the block. Those define sample data for development and can be later redefined before reusing the block. This functionality is not yet available in the data exploration environment discussed in Section 6.4, but we intend to add it in the future.

8 RELATED AND FUTURE WORK

This paper approaches the problem of live coding environments from a theoretical programming language perspective with a special focus on tooling for data science. Hence, the related and future work spans numerous areas.

Design and human-computer interaction. TODO: Victor [2012a]

From a design perspective, the idea of live programming environments has been popularised by Bret Victor [Victor 2012b]. Active research on novel forms of interaction happens in areas such as live coded music [Aaron and Blackwell 2013; Roberts et al. 2015]. The idea of live previews can be extended to direct manipulation [Shneiderman 1981]. The Gamma provides limited support for directly manipulating data (Section ??), but we intend to explore this direction further.

Data science tooling. An essential tool in data science is REPL (read-eval-print-loop) [Findler et al. 2002], which is now widely available. This has been integrated with rich graphical outputs in tools such as Jupyter notebooks [Kluyver et al. 2016; Pérez and Granger 2007], but such previews are updated using an explicit command. Integrating our work with Jupyter to provide instant live previews for R or Python would be an interesting extension of the presented work.

Live coding and live previews. Live previews have been implemented in LightTable [Granger 2012] and, more recently, in editors such as Chrome Developer Tools, but neither presents a simple description of their inner workings. An issue that attracts much attention is keeping state during code edits [Burckhardt et al. 2013; McDirmid 2007]. This would be an interesting problem if we extended our work to event-based reactive programming.

Structured editing. An alternative approach to ours is to avoid using text editors. Structured editors [Szwilius and Neal 1996] allow the user to edit the AST and could, in principle, recompute previews based on the performed operations, or preview evaluation as in interactive functional programming [Perera 2013]. A promising direction is using bi-directional lambda calculus [Omar et al. 2017]. Finally, abandoning text also enables building richer, more human-centric abstractions as illustrated by Subtext [Edwards 2005]. Our current focus, however, remains on text-based editors.

Dependency analysis. Our use of dependency graphs [Kuck et al. 1981] is first-order. Building dependency graphs involving function calls using modern compiler methods [Kennedy and Allen 2001] or program slicing [Weiser 1981] would allow us to deduce possible inputs for functions and use those for previews rather than changing the language as suggested in Section 7.2. This direction is worth considering, but it requires more empirical usability testing.

Semantics and partial evaluation. The evaluation of previews can be seen as a form of partial evaluation [Consel and Danvy 1993], done in a way that allows reuse of results. This can be done implicitly or explicitly in the form of multi-stage programming [Taha and Sheard 2000]. Both can provide useful perspective for formally analysing how previews are evaluated. Semantically, the evaluation of previews can be seen as a modality [Fairtlough et al. 2001] and delayed previews are linked to contextual modal type theory [Nanevski et al. 2008], which, in turn, can be understood in terms of comonads [Gabbay and Nanevski 2013]. This provides an intriguing direction for rigorous analysis of the presented system.

9 SUMMARY

We present The Gamma, a live coding environment for data exploration. The environment bridges the gap between spreadsheets and scripting – live previews give users rapid feedback, while the final result is a fully reproducible script.

In this paper, we focus on the challenge of efficiently providing live previews and type checking code during editing in a free-form text editor. This is a challenge, because users can perform arbitrary text transformations and we cannot recompute previews after each edit.

The key trick is to separate the process into fast *binding phase*, which constructs a dependency graph and slower *evaluation phase* and *type checking phase* that can cache results, using the nodes from the dependency graph created during binding as keys. This makes it possible to quickly

parse updated code, reconstruct dependency graph and compute preview using previous, partially evaluated, results.

We describe our approach formally, which serves two purposes. First, we aim to provide easy to use foundations for the growing and important trend of text-based live coding environments. Second, we explore the properties of our system and prove that our method does not recompute previews in a number of common cases and, at the same time, the optimisation still produces correct previews.

REFERENCES

- Samuel Aaron and Alan F Blackwell. 2013. From sonic Pi to overtone: creative musical experiences with domain-specific and functional languages. In *Proceedings of the first ACM SIGPLAN workshop on Functional art, music, modeling & design*. ACM, 35–46.
- Umut A. Acar. 2005. *Self-adjusting Computation*. Ph.D. Dissertation. Pittsburgh, PA, USA. AAI3166271.
- J Allaire, Joe Cheng, Yihui Xie, Jonathan McPherson, Winston Chang, Jeff Allen, Hadley Wickham, Aron Atkins, Rob Hyndman, and R Arslan. 2016. rmarkdown: Dynamic Documents for R. *R package version 1* (2016), 9010.
- David Blood. 2018. Recycling is broken – notebooks. Retrieved February 1, 2019 from <https://github.com/ft-interactive/recycling-is-broken-notebooks>
- Sebastian Burckhardt, Manuel Fahndrich, Peli de Halleux, Sean McDirmid, Michal Moskal, and Nikolai Tillmann. 2013. It's Alive! Continuous Feedback in UI Programming. In *PLDI*. ACM SIGPLAN.
- David Raymond Christiansen. 2013. Dependent type providers. In *Proceedings of the 9th ACM SIGPLAN workshop on Generic programming*. ACM, 25–34.
- Charles Consel and Olivier Danvy. 1993. Tutorial notes on partial evaluation. In *Proceedings of the 20th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM, 493–501.
- Andrew Crotty, Alex Galakatos, Emanuel Zraggen, Carsten Binnig, and Tim Kraska. 2015. Vizdom: Interactive Analytics Through Pen and Touch. *Proc. VLDB Endow.* 8, 12 (Aug. 2015), 2024–2027. <https://doi.org/10.14778/2824032.2824127>
- Jonathan Edwards. 2005. Subtext: uncovering the simplicity of programming. *ACM SIGPLAN Notices* 40, 10 (2005), 505–518.
- Matt Fairtlough, Michael Mendler, and Eugenio Moggi. 2001. Special issue: Modalities in type theory. *Mathematical Structures in Computer Science* 11, 4 (2001), 507–509.
- Robert Bruce Findler, John Clements, Cormac Flanagan, Matthew Flatt, Shriram Krishnamurthi, Paul Steckler, and Matthias Felleisen. 2002. DrScheme: A programming environment for Scheme. *Journal of functional programming* 12, 2 (2002), 159–182.
- Murdoch J Gabbay and Aleksandar Nanevski. 2013. Denotation of contextual modal type theory (CMTT): Syntax and meta-programming. *Journal of Applied Logic* 11, 1 (2013), 1–29.
- Marco Gaboardi, Shin-ya Katsumata, Dominic A Orchard, Flavien Breuvert, and Tarmo Uustalu. 2016. Combining effects and coeffects via grading.. In *ICFP*. 476–489.
- Chris Granger. 2012. LightTable: A new IDE concept. <http://www.chris-granger.com/2012/04/12/light-table-a-new-ide-concept/> (2012).
- Jonathan Gray, Lucy Chambers, and Liliana Bounegru. 2012. *The data journalism handbook: how journalists can use data to improve the news*. "O'Reilly Media, Inc".
- P Guo. 2013. Data science workflow: Overview and challenges. *blog CACM, Communications of the ACM* (2013).
- Matthew A. Hammer, Joshua Dunfield, Kyle Headley, Nicholas Labich, Jeffrey S. Foster, Michael Hicks, and David Van Horn. 2015. Incremental Computation with Names. In *Proceedings of the 2015 ACM SIGPLAN International Conference on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA 2015)*. ACM, New York, NY, USA, 748–766. <https://doi.org/10.1145/2814270.2814305>
- J. M. Hellerstein, R. Avnur, A. Chou, C. Hidber, C. Olston, V. Raman, T. Roth, and P. J. Haas. 1999. Interactive data analysis: the Control project. *Computer* 32, 8 (Aug 1999), 51–59. <https://doi.org/10.1109/2.781635>
- Leslie Hook and John Reed. 2018. Why the world's recycling system stopped working. Retrieved February 1, 2019 from <https://www.ft.com/content/360e2524-d71a-11e8-a854-33d6f82e62f8>
- Sean Kandel, Andreas Paepcke, Joseph Hellerstein, and Jeffrey Heer. 2011. Wrangler: Interactive Visual Specification of Data Transformation Scripts. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '11)*. ACM, 3363–3372. <https://doi.org/10.1145/1978942.1979444>
- Ken Kennedy and John R Allen. 2001. *Optimizing compilers for modern architectures: a dependence-based approach*. Morgan Kaufmann Publishers Inc.
- Thomas Kluyver, Benjamin Ragan-Kelley, Fernando Pérez, Brian E Granger, Matthias Bussonnier, Jonathan Frederic, Kyle Kelley, Jessica B Hamrick, Jason Grout, Sylvain Corlay, et al. 2016. Jupyter Notebooks-a publishing format for reproducible

- computational workflows.. In *ELPUB*, 87–90.
- David Koop and Jay Patel. 2017. Dataflow Notebooks: Encoding and Tracking Dependencies of Cells. In *9th USENIX Workshop on the Theory and Practice of Provenance, TaPP 2017, Seattle, WA, USA, June 23, 2017*.
- David J Kuck, Robert H Kuhn, David A Padua, Bruce Leasure, and Michael Wolfe. 1981. Dependence graphs and compiler optimizations. In *Proceedings of the 8th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM, 207–218.
- Eyal Lotem and Yair Chuchem. 2018. Lamdu Project. <https://github.com/lamdu/lamdu> (2018).
- Sean McDirmid. 2007. Living it up with a live programming language. In *ACM SIGPLAN Notices*, Vol. 42. ACM, 623–638.
- Wes McKinney. 2012. *Python for data analysis: Data wrangling with Pandas, NumPy, and IPython*. O'Reilly Media, Inc.
- Robin Milner. 1978. A theory of type polymorphism in programming. *J. Comput. System Sci.* 17, 3 (1978), 348 – 375. [https://doi.org/10.1016/0022-0000\(78\)90014-4](https://doi.org/10.1016/0022-0000(78)90014-4)
- Alan Mycroft, Dominic Orchard, and Tomas Petricek. 2016. Effect systems revisited—control-flow algebra and semantics. In *Semantics, Logics, and Calculi*. Springer, 1–32.
- Aleksandar Nanevski, Frank Pfenning, and Brigitte Pientka. 2008. Contextual modal type theory. *ACM Transactions on Computational Logic (TOCL)* 9, 3 (2008), 23.
- Cyrus Omar, Ian Voysey, Ravi Chugh, and Matthew A. Hammer. 2019. Live Functional Programming with Typed Holes. *PACMPL* 3, POPL (2019). <https://doi.org/10.1145/3290327>
- Cyrus Omar, Ian Voysey, Michael Hilton, Jonathan Aldrich, and Matthew A Hammer. 2017. Hazelnut: a bidirectionally typed structure editor calculus. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages*. ACM, 86–99.
- Roland Perera. 2013. *Interactive functional programming*. Ph.D. Dissertation. University of Birmingham.
- Fernando Pérez and Brian E Granger. 2007. IPython: a system for interactive scientific computing. *Computing in Science & Engineering* 9, 3 (2007).
- Tomas Petricek. 2017a. *Context-aware programming languages*. Ph.D. Dissertation. University of Cambridge, Computer Laboratory.
- Tomas Petricek. 2017b. Data Exploration through Dot-driven Development. In *LIPICs-Leibniz International Proceedings in Informatics*, Vol. 74. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik.
- Tomas Petricek, James Geddes, and Charles A. Sutton. 2018. Wrattler: Reproducible, live and polyglot notebooks. In *10th USENIX Workshop on the Theory and Practice of Provenance, TaPP 2018, London, UK, July 11-12, 2018*.
- Tomas Petricek, Gustavo Guerra, and Don Syme. 2016. Types from data: Making structured data first-class citizens in F. *ACM SIGPLAN Notices* 51, 6 (2016), 477–490.
- Charles Roberts, Matthew Wright, and JoAnn Kuchera-Morin. 2015. Beyond editing: extended interaction with textual code fragment. In *NIME*. 126–131.
- Peter Sestoft. 2012. *Spreadsheet technology*. Technical Report. Citeseer.
- Ben Shneiderman. 1981. Direct manipulation: A step beyond programming languages. In *ACM SIGSOC Bulletin*, Vol. 13. ACM, 143.
- Don Syme. 2006. Leveraging, net meta-programming components from F#: integrated queries and interoperable heterogeneous execution. In *Proceedings of the 2006 workshop on ML*. ACM, 43–54.
- Donald Syme, Keith Battocchi, Kenji Takeda, Donna Malayeri, and Tomas Petricek. 2013. Themes in information-rich functional programming for internet-scale data sources. In *Proceedings of the 2013 workshop on Data driven functional programming*. ACM, 1–4.
- Gerd Szwillus and Lisa Neal. 1996. *Structure-based editors and environments*. Academic Press, Inc.
- Walid Taha and Tim Sheard. 2000. MetaML and multi-stage programming with explicit annotations. *Theoretical computer science* 248, 1 (2000), 211–242.
- Bret Victor. 2012a. Inventing on Principle. <http://worrydream.com/InventingOnPrinciple> (2012).
- Bret Victor. 2012b. Learnable programming: Designing a programming system for understanding programs. <http://worrydream.com/LearnableProgramming> (2012).
- Mark Weiser. 1981. Program slicing. In *Proceedings of the 5th international conference on Software engineering*. IEEE Press, 439–449.
- Richard Wesley, Matthew Eldridge, and Pawel T. Terlecki. 2011. An Analytic Data Engine for Visualization in Tableau. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data (SIGMOD '11)*. ACM, 1185–1194. <https://doi.org/10.1145/1989323.1989449>

A APPENDIX

A.1 Normalization

Theorem 10 (Normalization). *For all p , there exists n and o_1, \dots, o_n such that $p \rightsquigarrow^* o_1; \dots; o_n$.*

PROOF. We define size of a program in data exploration calculus as follows:

$$\begin{aligned}
 \text{size}(c_1; \dots; c_n) &= 1 + \sum_{i=1}^n \text{size}(c_i) \\
 \text{size}(\text{let } x = t) &= 1 + \text{size}(t) \\
 \text{size}(e_0.m(e_1, \dots, e_n)) &= 1 + \sum_{i=0}^n \text{size}(e_i) \\
 \text{size}(\lambda x \rightarrow e) &= 1 + \text{size}(e) \\
 \text{size}(o) = \text{size}(x) &= 1
 \end{aligned} \tag{1}$$

The property holds because, first, both (let) and (external) decrease the size of the program and, second, a program is either fully evaluatevaluated, i.e. $o_1; \dots; o_n$ for some n or, it can be reduced using one of the reduction rules. \square

A.2 Term preview correctness

Theorem 11 (Term preview correctness). *Given a term t that has no free variables, together with a lookup table Δ obtained from any sequence of programs using bind-prog (Figure 6) and update (Figure 7), then let $\mathbf{v}, (V, E) = \text{bind-expr}_{\emptyset, \Delta}(t)$. If $\mathbf{v} \Downarrow p$ over a graph (V, E) then $p = o$ for some value o and $t \rightsquigarrow^* o$.*

PROOF. First note that, when combining recursively constructed sub-graphs, the bind-expr function adds new nodes and edges leading from those new nodes. Therefore, an evaluation using \Downarrow over a sub-graph will also be valid over the new graph – the newly added nodes and edges do not introduce non-determinism to the rules given in Figure 8.

We prove a more general property showing that for any e , its binding $\mathbf{v}, (V, E) = \text{bind-expr}_{\emptyset, \Delta}(e)$ and any evaluation context C such that $C[e] \rightsquigarrow o$ for some o , one of the following holds:

- If $FV(e) = \emptyset$ then $\mathbf{v} \Downarrow p$ for some p and $C[p] \rightsquigarrow o$
- If $FV(e) \neq \emptyset$ then $\mathbf{v} \Downarrow \llbracket e_p \rrbracket_{FV(e)}$ for some e_p and $C[e_p] \rightsquigarrow o$

In the first case, p is a value, but it is not always the case that $e \rightsquigarrow^* p$, because p may be lambda function and preview evaluation may reduce sub-expression in the body of the function. Using a context C in which the value reduces to an object avoids this problem.

The proof of the theorem follows from the more general property. Using a context $C[-] = -$, the term t reduces $t \rightsquigarrow^* t' \rightsquigarrow_e o$ for some o and the preview p is a value o because $C[p] = p = o$. The proof is by induction over the binding process, which follows the structure of the expression:

(1) $\text{bind-expr}_{\Gamma, \Delta}(e_0.m(e_1, \dots, e_n))$ – Here $e = e_0.m(e_1, \dots, e_n)$, \mathbf{v}_i are graph nodes obtained by induction for expressions e_i and $\{(v, \mathbf{v}_0, \text{arg}(0)), \dots, (v, \mathbf{v}_n, \text{arg}(n))\} \subseteq E$. From lookup inversion Lemma 4, $\mathbf{v} = \text{mem}(m, s)$ for some s .

If $FV(e) = \emptyset$, then $\mathbf{v}_i \Downarrow p_i$ for $i \in 0 \dots n$ and $\mathbf{v} \Downarrow p$ using (mem-val) such that $p_0.m(p_1, \dots, p_n) \rightsquigarrow p$. From induction hypothesis and *compositionality* of external libraries (Definition 2), it holds that for any C such that $C[e_0.m(e_1, \dots, e_n)] \rightsquigarrow o$ for some o then also $C[p_0.m(p_1, \dots, p_n)] \rightsquigarrow C[p] \rightsquigarrow o$.

If $FV(e) \neq \emptyset$, then $\mathbf{v}_i \Downarrow_{\text{lift}} \llbracket e'_i \rrbracket$ for $i \in 0 \dots n$ and $\mathbf{v} \Downarrow \llbracket e'_0.m(e'_1, \dots, e'_n) \rrbracket_{FV(e)}$ using (mem-expr). From induction hypothesis and *compositionality* of external libraries (Definition 2), it holds that for any C such that $C[e_0.m(e_1, \dots, e_n)] \rightsquigarrow o$ for some o then also $C[e'_0.m(e'_1, \dots, e'_n)] \rightsquigarrow o$.

(2) $\text{bind-expr}_{\Gamma, \Delta}(e_0.m(e_1, \dots, e_n))$ – This case is similar to (1), except that the fact that $\mathbf{v} = \text{mem}(m, s)$ holds by construction, rather than using Lemma 4.

(3) $\text{bind-expr}_{\Gamma, \Delta}(\lambda x \rightarrow e_b)$ – Here $e = \lambda x \rightarrow e_b$, v_b is the graph node obtained by induction for the expression e_b and $(v, v_b, \text{body}) \in E$. From lookup inversion Lemma 4, $v = \text{fun}(x, s)$ for some s . The evaluation can use one of three rules:

- i. If $FV(e) = \emptyset$ then $v_b \Downarrow p_b$ for some p_b and $v \Downarrow \lambda x \rightarrow p_b$ using (fun-val). Let $e'_b = p_b$.
- ii. If $FV(e_b) = \{x\}$ then $v_b \Downarrow \llbracket e'_b \rrbracket_x$ for some e'_b and $v \Downarrow \lambda x \rightarrow e'_b$ using (fun-bind).
- iii. Otherwise, $v_b \Downarrow \llbracket e'_b \rrbracket_{x, \Gamma}$ for some e'_b and $v \Downarrow \llbracket \lambda x \rightarrow e'_b \rrbracket_{\Gamma}$ using (fun-expr).

For i.) and ii.) we show that a.) is the case; for iii.) we show that b.) is the case; that is for any C , if $C[\lambda x \rightarrow e_b] \rightsquigarrow o$ then also $C[\lambda x \rightarrow e'_b] \rightsquigarrow o$. For a given C , let $C'[-] = C[\lambda x \rightarrow -]$ and use the induction hypothesis, i.e. if $C'[e_b] \rightsquigarrow o$ for some o then also $C'[e'_b] \rightsquigarrow o$.

(4) $\text{bind-expr}_{\Gamma, \Delta}(\lambda x \rightarrow e)$ – This case is similar to (3), except that the fact that $v = \text{fun}(x, s)$ holds by construction, rather than using Lemma 4.

(5) $\text{bind-expr}_{\Gamma, \Delta}(o)$ – In this case $e = o$ and $v = \text{val}(o)$ and $\text{val}(o) \Downarrow o$ using (val) and so the case a.) trivially holds.

(6) $\text{bind-expr}_{\Gamma, \Delta}(x)$ – The initial Γ is empty, so x must have been added to Γ by case (3) or (4). Hence, $v = \text{var}(x)$, $v \Downarrow \llbracket x \rrbracket_x$ using (var) and so $e_p = e = x$ and the case b.) trivially holds.

□

B APPENIDX - TYPE CHECKING

tbd